

1.5 Functions

In this section, we explore the creation of and use of **functions** in **Python**. As we did in Section 1.2.2, we draw a distinction between **functions** and **methods**. We use the general term **function** to describe a traditional, stateless **function** that is invoked without the context of a particular class or an instance of that class, such as `sorted(data)`. We use the more specific term **method** to describe a member **function** that is invoked upon a specific object using an object-oriented message passing syntax, such as `data.sort()`. In this section, we only consider pure **functions**; methods will be explored with more general object-oriented principles in Chapter 2.

We begin with an example to demonstrate the syntax for defining **functions** in **Python**. The following **function** counts the number of occurrences of a given target value within any form of iterable data set.

```
def count(data, target):
    n = 0
    for item in data:
        if item == target:                # found a match
            n += 1
    return n
```

The first line, beginning with the keyword **def**, serves as the **function**'s **signature**. This establishes a new identifier as the name of the **function** (`count`, in this example), and it establishes the number of parameters that it expects, as well as names identifying those parameters (`data` and `target`, in this example). Unlike Java and C++, **Python** is a dynamically typed language, and therefore a **Python** signature does not designate the types of those parameters, nor the type (if any) of a return value. Those expectations should be stated in the **function**'s documentation (see Section 2.2.3) and can be enforced within the body of the **function**, but misuse of a **function** will only be detected at run-time.

The remainder of the **function** definition is known as the **body** of the **function**. As is the case with control structures in **Python**, the body of a **function** is typically expressed as an indented block of code. Each time a **function** is called, **Python** creates a dedicated **activation record** that stores information relevant to the current call. This activation record includes what is known as a **namespace** (see Section 1.10) to manage all identifiers that have **local scope** within the current call. The namespace includes the **function**'s parameters and any other identifiers that are defined locally within the body of the **function**. An identifier in the local scope of the **function** caller has no relation to any identifier with the same name in the caller's scope (although identifiers in different scopes may be aliases to the same object). In our first example, the identifier `n` has scope that is local to the **function** call, as does the identifier `item`, which is established as the loop variable.

1.12 Exercises

For help with exercises, please visit the site, www.wiley.com/college/goodrich.

Reinforcement

- R-1.1** Write a short **Python function**, `is_multiple(n, m)`, that takes two integer values and returns `True` if n is a multiple of m , that is, $n = mi$ for some integer i , and `False` otherwise.
- R-1.2** Write a short **Python function**, `is_even(k)`, that takes an integer value and returns `True` if k is even, and **False** otherwise. However, your **function** cannot use the multiplication, modulo, or division operators.
- R-1.3** Write a short **Python function**, `minmax(data)`, that takes a sequence of one or more numbers, and returns the smallest and largest numbers, in the form of a tuple of length two. Do not use the built-in **functions** `min` or `max` in implementing your solution.
- R-1.4** Write a short **Python function** that takes a positive integer n and returns the sum of the squares of all the positive integers smaller than n .
- R-1.5** Give a single command that computes the sum from Exercise R-1.4, relying on **Python's** comprehension syntax and the built-in `sum` **function**.
- R-1.6** Write a short **Python function** that takes a positive integer n and returns the sum of the squares of all the odd positive integers smaller than n .
- R-1.7** Give a single command that computes the sum from Exercise R-1.6, relying on **Python's** comprehension syntax and the built-in `sum` **function**.
- R-1.8** **Python** allows negative integers to be used as indices into a sequence, such as a string. If string s has length n , and expression $s[k]$ is used for index $-n \leq k < 0$, what is the equivalent index $j \geq 0$ such that $s[j]$ references the same element?
- R-1.9** What parameters should be sent to the `range` constructor, to produce a range with values 50, 60, 70, 80?
- R-1.10** What parameters should be sent to the `range` constructor, to produce a range with values 8, 6, 4, 2, 0, -2, -4, -6, -8?
- R-1.11** Demonstrate how to use **Python's** list comprehension syntax to produce the list `[1, 2, 4, 8, 16, 32, 64, 128, 256]`.
- R-1.12** **Python's** random module includes a **function** `choice(data)` that returns a random element from a non-empty sequence. The random module includes a more basic **function** `randrange`, with parameterization similar to the built-in `range` **function**, that return a random choice from the given range. Using only the `randrange` **function**, implement your own version of the `choice` **function**.

15.1.3 Additional Memory Used by the Python Interpreter

We have discussed, in Section 15.1.1, how the Python interpreter allocates memory for objects within a memory heap. However, this is not the only memory that is used when executing a Python program. In this section, we discuss some other important uses of memory.

The Run-Time Call Stack

Stacks have a most important application to the run-time environment of Python programs. A running Python program has a private stack, known as the *call stack* or *Python interpreter stack*, that is used to keep track of the nested sequence of currently active (that is, nonterminated) invocations of functions. Each entry of the stack is a structure known as an *activation record* or *frame*, storing important information about an invocation of a function.

At the top of the call stack is the activation record of the *running call*, that is, the function activation that currently has control of the execution. The remaining elements of the stack are activation records of the *suspended calls*, that is, functions that have invoked another function and are currently waiting for that other function to return control when it terminates. The order of the elements in the stack corresponds to the chain of invocations of the currently active functions. When a new function is called, an activation record for that call is pushed onto the stack. When it terminates, its activation record is popped from the stack and the Python interpreter resumes the processing of the previously suspended call.

Each activation record includes a dictionary representing the local namespace for the function call. (See Sections 1.10 and 2.5 for further discussion of namespaces). The namespace maps identifiers, which serve as parameters and local variables, to object values, although the objects being referenced still reside in the memory heap. The activation record for a function call also includes a reference to the function definition itself, and a special variable, known as the *program counter*, to maintain the address of the statement within the function that is currently executing. When one function returns control to another, the stored program counter for the suspended function allows the interpreter to properly continue execution of that function.

Implementing Recursion

One of the benefits of using a stack to implement the nesting of function calls is that it allows programs to use *recursion*. That is, it allows a function to call itself, as discussed in Chapter 4. We implicitly described the concept of the call stack and the use of activation records within our portrayal of *recursion traces* in

When an identifier is indicated in a command, **Python** searches a series of namespaces in the process of name resolution. First, the most locally enclosing scope is searched for a given name. If not found there, the next outer scope is searched, and so on. We will continue our examination of namespaces, in Section 2.5, when discussing **Python**'s treatment of object-orientation. We will see that each object has its own namespace to store its attributes, and that classes each have a namespace as well.

First-Class Objects

In the terminology of programming languages, *first-class objects* are instances of a type that can be assigned to an identifier, passed as a parameter, or returned by a **function**. All of the data types we introduced in Section 1.2.3, such as `int` and `list`, are clearly first-class types in **Python**. In **Python**, **functions** and **classes** are also treated as first-class objects. For example, we could write the following:

```
scream = print      # assign name 'scream' to the function denoted as 'print'
scream('Hello')    # call that function
```

In this case, we have not created a new **function**, we have simply defined `scream` as an alias for the existing `print` **function**. While there is little motivation for precisely this example, it demonstrates the mechanism that is used by **Python** to allow one **function** to be passed as a parameter to another. On page 28, we noted that the built-in **function**, `max`, accepts an optional keyword parameter to specify a non-default order when computing a maximum. For example, a caller can use the syntax, `max(a, b, key=abs)`, to determine which value has the larger absolute value. Within the body of that **function**, the formal parameter, `key`, is an identifier that will be assigned to the actual parameter, `abs`.

In terms of namespaces, an assignment such as `scream = print`, introduces the identifier, `scream`, into the current namespace, with its value being the object that represents the built-in **function**, `print`. The same mechanism is applied when a user-defined **function** is declared. For example, our `count` **function** from Section 1.5 beings with the following syntax:

```
def count(data, target):
    ...
```

Such a declaration introduces the identifier, `count`, into the current namespace, with the value being a **function** instance representing its implementation. In similar fashion, the name of a newly defined class is associated with a representation of that class as its value. (Class definitions will be introduced in the next chapter.)

Creativity

- C-1.13** Write a pseudo-code description of a **function** that reverses a list of n integers, so that the numbers are listed in the opposite order than they were before, and compare this method to an equivalent **Python function** for doing the same thing.
- C-1.14** Write a short **Python function** that takes a sequence of integer values and determines if there is a distinct pair of numbers in the sequence whose product is odd.
- C-1.15** Write a **Python function** that takes a sequence of numbers and determines if all the numbers are different from each other (that is, they are distinct).
- C-1.16** In our implementation of the **scale function** (page 25), the body of the loop executes the command `data[j] *= factor`. We have discussed that numeric types are immutable, and that use of the `*=` operator in this context causes the creation of a new instance (not the mutation of an existing instance). How is it still possible, then, that our implementation of **scale** changes the actual parameter sent by the caller?
- C-1.17** Had we implemented the **scale function** (page 25) as follows, does it work properly?

```
def scale(data, factor):  
    for val in data:  
        val *= factor
```

Explain why or why not.

- C-1.18** Demonstrate how to use **Python's** list comprehension syntax to produce the list `[0, 2, 6, 12, 20, 30, 42, 56, 72, 90]`.
- C-1.19** Demonstrate how to use **Python's** list comprehension syntax to produce the list `['a', 'b', 'c', ..., 'z']`, but without having to type all 26 such characters literally.
- C-1.20** **Python's** random module includes a **function** `shuffle(data)` that accepts a list of elements and randomly reorders the elements so that each possible order occurs with equal probability. The random module includes a more basic **function** `randint(a, b)` that returns a uniformly random integer from a to b (including both endpoints). Using only the `randint function`, implement your own version of the **shuffle function**.
- C-1.21** Write a **Python** program that repeatedly reads lines from standard input until an `EOFError` is raised, and then outputs those lines in reverse order (a user can indicate end of input by typing `ctrl-D`).

12.6 Python's Built-In Sorting Functions

Python provides two built-in ways to sort data. The first is the sort method of the list class. As an example, suppose that we define the following list:

```
colors = ['red', 'green', 'blue', 'cyan', 'magenta', 'yellow']
```

That method has the effect of reordering the elements of the list into order, as defined by the natural meaning of the < operator for those elements. In the above example, within elements that are strings, the natural order is defined alphabetically. Therefore, after a call to colors.sort(), the order of the list would become:

```
['blue', 'cyan', 'green', 'magenta', 'red', 'yellow']
```

Python also supports a built-in function, named sorted, that can be used to produce a new ordered list containing the elements of any existing iterable container. Going back to our original example, the syntax sorted(colors) would return a new list of those colors, in alphabetical order, while leaving the contents of the original list unchanged. This second form is more general because it can be applied to any iterable object as a parameter; for example, sorted('green') returns ['e', 'e', 'g', 'n', 'r'].

12.6.1 Sorting According to a Key Function

There are many situations in which we wish to sort a list of elements, but according to some order other than the natural order defined by the < operator. For example, we might wish to sort a list of strings from shortest to longest (rather than alphabetically). Both of Python's built-in sort functions allow a caller to control the notion of order that is used when sorting. This is accomplished by providing, as an optional keyword parameter, a reference to a secondary function that computes a *key* for each element of the primary sequence; then the primary elements are sorted based on the natural order of their keys. (See pages 27 and 28 of Section 1.5.1 for a discussion of this technique in the context of the built-in min and max functions.)

A key function must be a one-parameter function that accepts an element as a parameter and returns a key. For example, we could use the built-in len function when sorting strings by length, as a call len(s) for string s returns its length. To sort our colors list based on length, we use the syntax colors.sort(key=len) to mutate the list or sorted(colors, key=len) to generate a new ordered list, while leaving the original alone. When sorted with the length function as a key, the contents are:

```
['red', 'blue', 'cyan', 'green', 'yellow', 'magenta']
```

These built-in functions also support a keyword parameter, reverse, that can be set to True to cause the sort order to be from largest to smallest.

This **function** does not use any explicit loops. Repetition is provided by the repeated recursive invocations of the **function**. There is no circularity in this definition, because each time the **function** is invoked, its argument is smaller by one, and when a base case is reached, no further recursive calls are made.

We illustrate the execution of a recursive **function** using a **recursion trace**. Each entry of the trace corresponds to a recursive call. Each new recursive **function** call is indicated by a downward arrow to a new invocation. When the **function** returns, an arrow showing this return is drawn and the return value may be indicated alongside this arrow. An example of such a trace for the factorial **function** is shown in Figure 4.1.

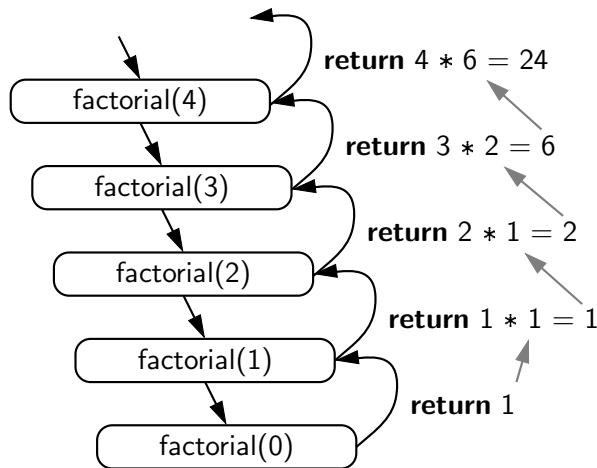


Figure 4.1: A recursion trace for the call `factorial(5)`.

A recursion trace closely mirrors the programming language's execution of the recursion. In **Python**, each time a **function** (recursive or otherwise) is called, a structure known as an **activation record** or **frame** is created to store information about the progress of that invocation of the **function**. This activation record includes a namespace for storing the **function** call's parameters and local variables (see Section 1.10 for a discussion of namespaces), and information about which command in the body of the **function** is currently executing.

When the execution of a **function** leads to a nested **function** call, the execution of the former call is suspended and its activation record stores the place in the source code at which the flow of control should continue upon return of the nested call. This process is used both in the standard case of one **function** calling a different **function**, or in the recursive case in which a **function** invokes itself. The key point is that there is a different activation record for each active call.

Return Statement

A **return** statement is used within the body of a **function** to indicate that the **function** should immediately cease execution, and that an expressed value should be returned to the caller. If a return statement is executed without an explicit argument, the `None` value is automatically returned. Likewise, `None` will be returned if the flow of control ever reaches the end of a **function** body without having executed a return statement. Often, a return statement will be the final command within the body of the **function**, as was the case in our earlier example of a `count` **function**. However, there can be multiple return statements in the same **function**, with conditional logic controlling which such command is executed, if any. As a further example, consider the following **function** that tests if a value exists in a sequence.

```
def contains(data, target):
    for item in target:
        if item == target:                # found a match
            return True
    return False
```

If the conditional within the loop body is ever satisfied, the `return True` statement is executed and the **function** immediately ends, with `True` designating that the target value was found. Conversely, if the `for` loop reaches its conclusion without ever finding the match, the final `return False` statement will be executed.

1.5.1 Information Passing

To be a successful programmer, one must have clear understanding of the mechanism in which a programming language passes information to and from a **function**. In the context of a **function** signature, the identifiers used to describe the expected parameters are known as *formal parameters*, and the objects sent by the caller when invoking the **function** are the *actual parameters*. Parameter passing in **Python** follows the semantics of the standard *assignment statement*. When a **function** is invoked, each identifier that serves as a formal parameter is assigned, in the **function**'s local scope, to the respective actual parameter that is provided by the caller of the **function**.

For example, consider the following call to our `count` **function** from page 23:

```
prizes = count(grades, 'A')
```

Just before the **function** body is executed, the actual parameters, `grades` and `'A'`, are implicitly assigned to the formal parameters, `data` and `target`, as follows:

```
data = grades
target = 'A'
```


These assignment statements establish identifier data as an alias for grades and target as a name for the string literal 'A'. (See Figure 1.7.)




Figure 1.7: A portrayal of parameter passing in Python, for the function call `count(grades, 'A')`. Identifiers `data` and `target` are formal parameters defined within the local scope of the `count` function.

The communication of a return value from the function back to the caller is similarly implemented as an assignment. Therefore, with our sample invocation of `prizes = count(grades, 'A')`, the identifier `prizes` in the caller's scope is assigned to the object that is identified as `n` in the return statement within our function body.

An advantage to Python's mechanism for passing information to and from a function is that objects are not copied. This ensures that the invocation of a function is efficient, even in a case where a parameter or return value is a complex object.

Mutable Parameters

Python's parameter passing model has additional implications when a parameter is a mutable object. Because the formal parameter is an alias for the actual parameter, the body of the function may interact with the object in ways that change its state. Considering again our sample invocation of the `count` function, if the body of the function executes the command `data.append('F')`, the new entry is added to the end of the list identified as `data` within the function, which is one and the same as the list known to the caller as `grades`. As an aside, we note that reassigning a new value to a formal parameter with a function body, such as by setting `data = []`, does not alter the actual parameter; such a reassignment simply breaks the alias.

Our hypothetical example of a `count` method that appends a new element to a list lacks common sense. There is no reason to expect such a behavior, and it would be quite a poor design to have such an unexpected effect on the parameter. There are, however, many legitimate cases in which a function may be designed (and clearly documented) to modify the state of a parameter. As a concrete example, we present the following implementation of a method named `scale` that's primary purpose is to multiply all entries of a numeric data set by a given factor.

```

def scale(data, factor):
    for j in range(len(data)):
        data[j] *= factor
  
```

By default, `max` operates based upon the natural order of elements according to the `<` operator for that type. But the maximum can be computed by comparing some other aspect of the elements. This is done by providing an auxiliary *function* that converts a natural element to some other value for the sake of comparison. For example, if we are interested in finding a numeric value with *magnitude* that is maximal (i.e., considering -35 to be larger than $+20$), we can use the calling syntax `max(a, b, key=abs)`. In this case, the built-in `abs` *function* is itself sent as the value associated with the keyword parameter `key`. (Functions are first-class objects in Python; see Section 1.10.) When `max` is called in this way, it will compare `abs(a)` to `abs(b)`, rather than `a` to `b`. The motivation for the keyword syntax as an alternate to positional arguments is important in the case of `max`. This *function* is polymorphic in the number of arguments, allowing a call such as `max(a,b,c,d)`; therefore, it is not possible to designate a key *function* as a traditional positional element. Sorting *functions* in Python also support a similar `key` parameter for indicating a nonstandard order. (We explore this further in Section 9.4 and in Section 12.6.1, when discussing sorting algorithms).

1.5.2 Python's Built-In Functions

Table 1.4 provides an overview of common *functions* that are automatically available in Python, including the previously discussed `abs`, `max`, and `range`. When choosing names for the parameters, we use identifiers `x`, `y`, `z` for arbitrary numeric types, `k` for an integer, and `a`, `b`, and `c` for arbitrary comparable types. We use the identifier, `iterable`, to represent an instance of any iterable type (e.g., `str`, `list`, `tuple`, `set`, `dict`); we will discuss iterators and iterable data types in Section 1.8. A sequence represents a more narrow category of indexable classes, including `str`, `list`, and `tuple`, but neither `set` nor `dict`. Most of the entries in Table 1.4 can be categorized according to their *functionality* as follows:

Input/Output: `print`, `input`, and `open` will be more fully explained in Section 1.6.

Character Encoding: `ord` and `chr` relate characters and their integer code points. For example, `ord('A')` is 65 and `chr(65)` is `'A'`.

Mathematics: `abs`, `divmod`, `pow`, `round`, and `sum` provide common mathematical *functionality*; an additional math module will be introduced in Section 1.11.

Ordering: `max` and `min` apply to any data type that supports a notion of comparison, or to any collection of such values. Likewise, `sorted` can be used to produce an ordered list of elements drawn from any existing collection.

Collections/Iterations: `range` generates a new sequence of numbers; `len` reports the length of any existing collection; *functions* `reversed`, `all`, `any`, and `map` operate on arbitrary iterations as well; `iter` and `next` provide a general framework for iteration through elements of a collection, and are discussed in Section 1.8.